Classical Machine Learning

Week 0

Plan

- Setting up your learning and programming environment

Getting started

- [Setting up your ML environment (Setup_NYU.ipynb)](#)
    - [Choosing an ML environment (Choosing_an_ML_Environment_NYU.ipynb)](#)
- [Quick intro to the tools (Getting_Started.ipynb)](#)

Week 1

Plan

- Motivate Machine Learning
- Introduce notation used throughout course
- Plan for initial lectures
    - *What*: Introduce, motivate a model
    - *How*: How to use a model: function signature, code (API)
    - *Why*: Mathematical basis -- enhance understanding and ability to improve results

- [Course Overview (Course_overview_NYU.ipynb)](Course_overview_NYU.ipynb)
- [Machine Learning: Overview (ML_Overview.ipynb)](ML_Overview.ipynb)
- [Intro to Classical ML (Intro_Classical_ML.ipynb)](Intro_Classical_ML.ipynb)

Using an AI Assistant

AI Assistants are often very good at coding.

But using one to just "get the answer" deprives you of a valuable tool

- you can ask the Assistant *why* it chose to do something
- keep on asking
- treat it as a private tutor !

[Learning about KNN using an Assistant as a private tutor (https://www.perplexity.ai/search/using-python-and-sklearn-pleas-407oe3uzTXu1i9xEHVR2MQ)](https://www.perplexity.ai/search/using-python-and-sklearn-pleas-407oe3uzTXu1i9xEHVR2MQ)

Week 2

Recap of last week

- [Summary of Intro to Supervised Machine Learning (Intro_to_Supervised_Learning_Summary.ipynb)](Intro_to_Supervised_Learning_Summary.ipynb)
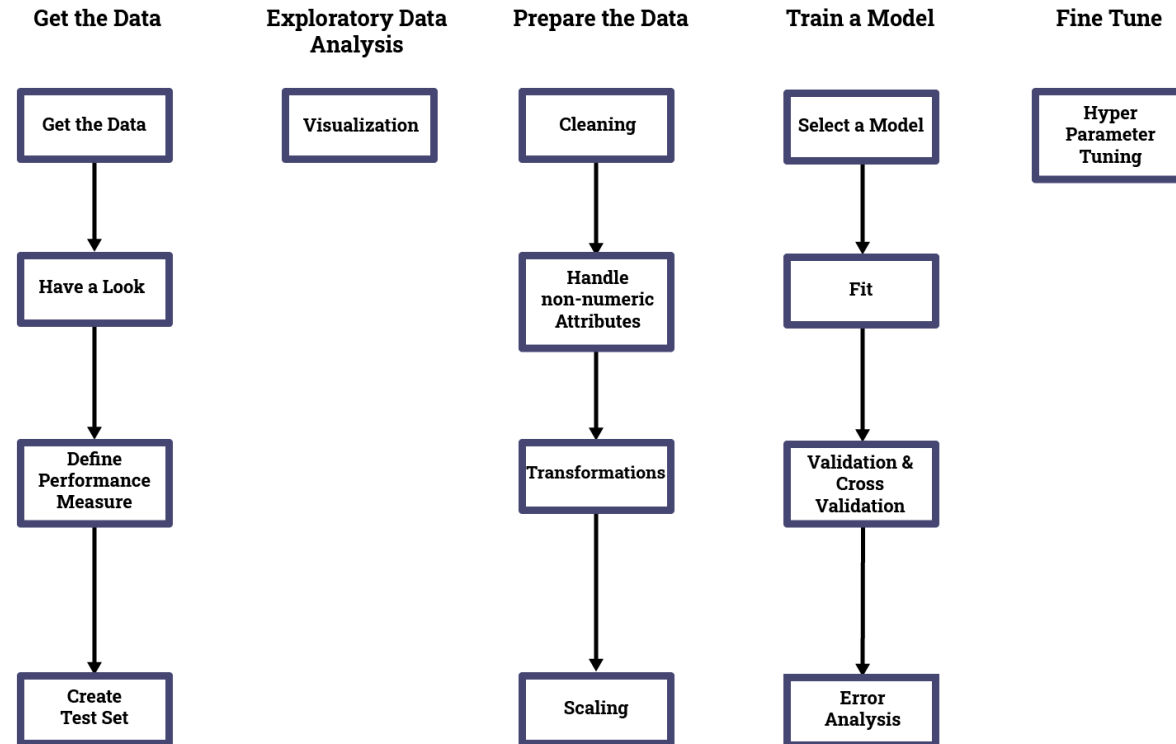
Plan

We will learn the Recipe for Machine Learning, a disciplined approach to solving problems in Machine Learning.

We will illustrate the Recipe while, at the same time, introducing a model for the Regression task: Linear Regression.

Our coverage of the Recipe will be rapid and shallow (we use an extremely simple example for illustration).

I highly recommend reviewing and understanding this [Geron notebook (external/handson-ml2/02_end_to_end_machine_learning_project.ipynb)](external/handson-ml2/02_end_to_end_machine_learning_project.ipynb) in order to acquire a more in-depth appreciation of the Recipe.

Get the Data	Exploratory Data Analysis	Prepare the Data	Train a Model	Fine Tune
Get the Data	Visualization	Cleaning	Select a Model	Hyper Parameter Tuning
Have a Look		Handle non-numeric Attributes	Fit	
Define Performance Measure		Transformations	Validation & Cross Validation	
Create Test Set		Scaling	Error Analysis	

Recipe, as illustrated by Linear Regression

[The Recipe for Machine Learning: Solving a Regression task (Recipe_via_Linear_Regression.ipynb)](Recipe_via_Linear_Regression.ipynb)

- A *process* for Machine Learning
  - Go through the methodical, multi-step process
    - Quick first pass, followed by Deeper Dives

Fitting a model: details

Recall: fitting a model (finding optimal value for the parameters) is found by minimizing a Loss function.

Let's examine a typical Loss function for Regression

- [Regression: Loss Function (Linear_Regression_Loss_Function.ipynb)](Linear_Regression_Loss_Function.ipynb)

Iterative training: when to stop

Increasing the number of parameters of a model improves in-sample fit (reduces Loss) but may compromise out-of-sample prediction (generalization).

We examine the issues of having too many/too few parameters.

- [When to stop iterating: Bias and Variance (Bias_and_Variance.ipynb)](Bias_and_Variance.ipynb)

Get the data: Fundamental Assumption of Machine Learning

- [Getting *good* training examples (Recipe_Training_data.ipynb)](Recipe_Training_data.ipynb)

Regression: final thoughts (for now)

- [Regression: coda (Regression_coda.ipynb)](Regression_coda.ipynb)

Deeper dives

- [Fine tuning techniques (Fine_tuning.ipynb)](#)

Using an AI Assistant

[Learning about Linear Regression using an Assistant as private tutor (https://www.perplexity.ai/search/using-python-sklearn-and-matpl-vTYv7oGdRO6upR5L5OSirg)](#)

Week 3

Recipe "Prepare the Data" step: Transformations

We covered this at end of Week 2

We recap the importance of adding *synthetic* features to our Linear Regression example

- and *preview* the *mechanical* process of creating these features via *Transformations*

Transformations

- [Prepare Data: Intro to Transformations (Prepare_data_Overview.ipynb)](#)

Validation

We skipped over this topic in Week 2; we cover it now

Our test dataset can be used only once, yet

- we have an iterative process for developing models
- each iteration requires a proxy for out of sample data to use in the Performance Metric

The solution: create a proxy for out of sample that is a *subset* of the training data.

- [Validation and Cross-Validation (Recipe_via_Linear_Regression.ipynb#Validation-and-Cross-Validation)](Recipe_via_Linear_Regression.ipynb#Validation-and-Cross-Validation)
- [Avoiding cheating in Cross-Validation (Prepare_data_Overview.ipynb#Using-pipelines-to-avoid-cheating-in-cross-validation)](Prepare_data_Overview.ipynb#Using-pipelines-to-avoid-cheating-in-cross-validation)

Classification Task

Plan

- We introduce a model for the Classification task: Logistic Regression
- How to deal with Categorical (non-numeric) variables
    - classification target
    - features

Classification intro

- [Classification: Overview (Classification_Overview.ipynb)](Classification_Overview.ipynb)
- [Classification and Categorical Variables (Classification_Notebook_Overview.ipynb)](Classification_Notebook_Overview.ipynb)
  - [linked notebook (Classification_and_Non_Numerical_Data.ipynb)](Classification_and_Non_Numerical_Data.ipynb)

Categorical variables (contained as subsections of Classification and Categorical Variables)

We examine the proper treatment of categorical variables (target or feature).

Along the way, we run into a subtle difficulty: the Dummy Variable Trap.

- [Classification and Categorical Variables: Categorical Variables (Classification_Notebook_Overview.ipynb#Categorical-variables)](Classification_Notebook_Overview.ipynb#Categorical-variables)

Week 4

Multinomial Classification

We generalize Binary Classification into classification into more than two classes.

- [Multinomial Classification (Multinomial_Classification.ipynb)](Multinomial_Classification.ipynb)

Classification and Categorical variables wrapup

- [Classification Loss Function (Classification_Loss_Function.ipynb)](Classification_Loss_Function.ipynb)
- [Baseline model for Classification (Classification_Baseline_Model.ipynb)](Classification_Baseline_Model.ipynb)
- [OHE issue: Dummy variable trap (Dummy_Variable_Trap.ipynb)](Dummy_Variable_Trap.ipynb)

Classification: final thoughts (for now)

- [Classification: coda (Classification_coda.ipynb)](Classification_coda.ipynb)

Plan

Good news

- You now know two main tasks in Supervised Learning
    - Regression, Classification
- You now know how to use virtually every model in `sklearn`
    - Consistent API
        - `fit`, `transform`, `predict`
- You survived the "sprint" to get you up and running with ML
- You know the *mechanical process* to implement transformations: Pipelines

Time to re-visit, in more depth, several important topics

Error Analysis

- We explain Error Analysis for the Classification Task, with a detailed example

- How Training Loss can be improved

- [Error Analysis (Error_Analysis_Overview.ipynb)](#)

    - [linked notebook (Error_Analysis.ipynb)](#)
        - Summary statistics
        - Conditional statistics
    - [Worked example (Error_Analysis_MNIST.ipynb)](#)

- [Loss Analysis: Using training loss to improve models (Training_Loss.ipynb)](#)

Imbalanced data

- [Imbalanced data (Imbalanced_Data.ipynb)](#)

Transformations: the "why"

Part of becoming a better Data Scientist is transforming raw features into more useful synthetic features.

In an earlier week, we presented the [mechanics (Prepare_data_Overview.ipynb)](#) (how to use `sklearn` to implement transformation Pipelines) of Transformations. This week, we focus on the necessity (the "why"): transforming raw data into something that tells a story

- [Becoming a successful Data Scientist (Becoming_a_successful_Data_Scientist.ipynb)](#)
- [Transformations: overview (Transformations_Overview.ipynb)](#)
    - linked notebooks:

Week 5

Transformations: the "why" (continued)

- [Transformations: overview (Transformations_Overview.ipynb#Scaling)](Transformations_Overview.ipynb#Scaling)
    - linked notebooks:
        - [Transformations: scaling (Transformations_Scaling.ipynb)](Transformations_Scaling.ipynb)
        - [Transformations: normalization (Transformations_Normalization.ipynb)](Transformations_Normalization.ipynb)
        - [Other Transformations (Transformations_Other.ipynb)](Transformations_Other.ipynb)

Loss functions: mathematical basis

Where do the Loss functions of Classical Machine Learning come from ? We take a brief mathematical detour into Loss functions.

- [Entropy, Cross Entropy, and KL Divergence (Entropy_Cross_Entropy_KL_Divergence.ipynb)](Entropy_Cross_Entropy_KL_Divergence.ipynb)
- [Loss functions: the math (Loss_functions.ipynb)](Loss_functions.ipynb)
    - Maximum likelihood
    - Preview: custom loss functions and Deep Learning

More models for classification

Plan

- More models: Decision Trees, Naive Bayes
    - Different flavor: more procedural, less mathematical
    - Decision Trees: a model with *non-linear* boundaries
- Ensembles
    - Bagging and Boosting
    - Random Forests

Decision Trees, Ensembles

- [Decision Trees: Overview (Decision_Trees_Overview.ipynb)](#)
- [Decision Trees (Decision_Trees_Notebook_Overview.ipynb)](#)
    - [linked notebook (Decision_Trees.ipynb)](#)
- [Trees, Forests, Ensembles (Ensembles.ipynb)](#)

Naive Bayes

Additional Deep Learning resources

Here are some resources that I have found very useful.

Some of them are very nitty-gritty, deep-in-the-weeds (even the "introductory" courses)

- For example: let's make believe PyTorch (or Keras/TensorFlow) didn't exists; let's invent Deep Learning without it !
    - You will gain a deeper appreciation and understanding by re-inventing that which you take for granted

[Andrej Karpathy course: Neural Networks, Zero to Hero (https://karpathy.ai/zero-to-hero.html)](https://karpathy.ai/zero-to-hero.html)

- PyTorch
- Introductory, but at a very deep level of understanding
    - you will get very deep into the weeds (hand-coding gradients !) but develop a deeper appreciation

fast.ai

`fast.ai` is a web-site with free courses from Jeremy Howard.

- PyTorch
- Introductory and courses "for coders"
- Same courses offered every few years, but sufficiently different so as to make it worthwhile to repeat the course !
  - [Practical Deep Learning (https://course.fast.ai/)](https://course.fast.ai/)
  - [Stable diffusion (https://course.fast.ai/Lessons/part2.html)](https://course.fast.ai/Lessons/part2.html)
    - Very detailed, nitty-gritty details (like Karpathy) that will give you a deeper appreciation

[Stefan Jansen: Machine Learning for Trading (https://github.com/stefan-jansen/machine-learning-for-trading)](https://github.com/stefan-jansen/machine-learning-for-trading)

An excellent github repo with notebooks

- using Deep Learning for trading
- Keras
- many notebooks are cleaner implementations of published models

Assignments

Your assignments should follow the [Assignment Guidelines (assignments/Assignment_Guidelines.ipynb)](assignments/Assignment_Guidelines.ipynb)

Regression

- Assignment notebook: [Using Machine Learning for Hedging (assignments/Regression%20task/Using_Machine_Learning_for_Hedging.ipynb)](assignments/Regression%20task/Using_Machine_Learning_for_Hedging.ipynb)
- Data
    - There is an archive file containing the data
    - You can find it
        - Under the course page: Content --> Data --> Assignments --> Regression task
        - You won't be able to view the file in the browser, but you **will** be able to Download it
    - You should unzip this archive into the *the same directory* as the assignment notebook
    - The end result is that the directory should contain
        - The assignment notebook and a helper file
        - A directory named `Data`

```
In [1]: print("Done")
```

Done