Neural Network

A Neural Network (Sequential architecture) is composed of

- sequence of Layers
  - layer $\ll$ transforms its input $\backslash \mathrm{y}_{(\ll -1)}$ to output $\backslash \mathrm{y}_{\backslash \mathrm{llp}}$
    - through a transformation: operation parameterized by weights $\backslash \mathrm{W}_{\backslash \mathrm{llp}}$

In [6]:
```
fig_tf_seq
```

Out[6]:

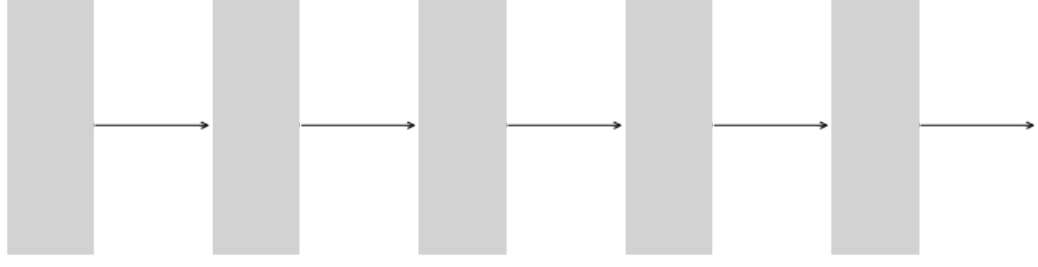

- initial layer 0 is Input layer:
  - outputs the network's inputs $\mathbf{x}$

$$\mathbf{y}_{(0)} = \mathbf{x}$$

- final layer L transforms its input $\mathbf{y}_{(L-1)}$ to prediction $\hat{\mathbf{y}}$

$$\hat{\mathbf{y}} = \mathbf{y}_{\mathrm{llp}}$$

  - transformation of layer L usually: Regression or Classification

In [7]:
```
fig_tf_test
```

Out[7]:

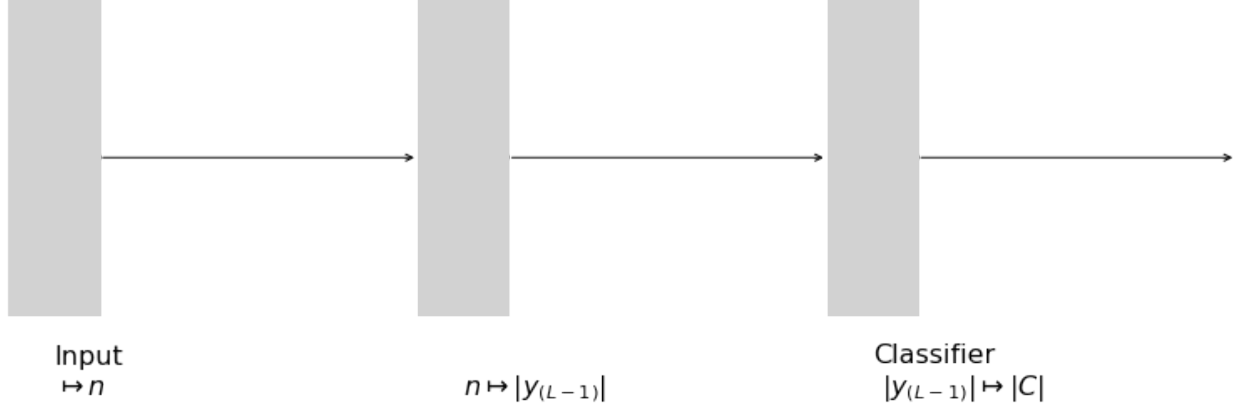

In the above diagram

- the central box represents a sequence of 1 or more layers
- a sub-network
- just for brevity

The Neural network thus computes a *function* from $\mathbf{x}$ to $\hat{\mathbf{y}}$.

The function *mimics* the training data

$$\langle \mathbf{X}, \mathbf{y} \rangle = [\mathbf{x}^{(i)}, \mathbf{y}^{(i)} | 1 \leq i \leq m]$$

where each *example* $\langle \mathbf{x}^{(i)}, \mathbf{y}^{(i)} \rangle$

- describes the mapping of the function on input $\mathbf{x}^{(i)}$ to output $\mathbf{y}^{(i)}$
- e.g., from input features $\mathbf{x}^{(i)}$ to
  - continuous value $\mathbf{y}^{(i)}$
  - discrete class $\mathbf{y}^{(i)}$
    - really: output is a probability vector over finite set C of discrete classes

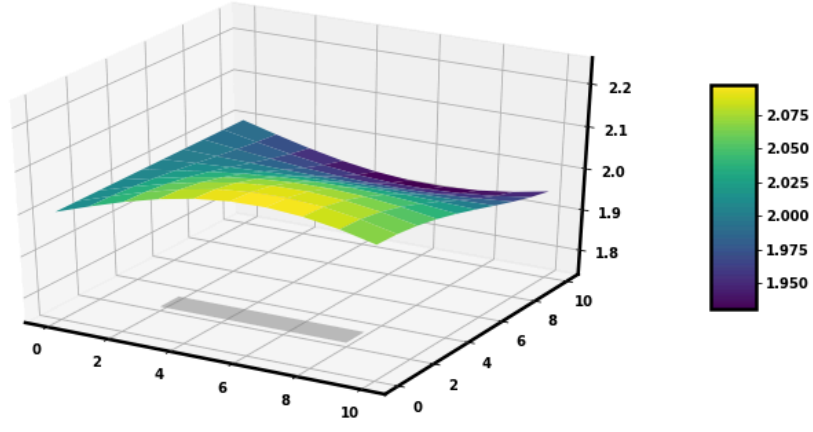
2.2
2.1
2.0
1.9
1.8
0
2
4
6
8
10
2.075
2.050
2.025
2.000
1.975
1.950

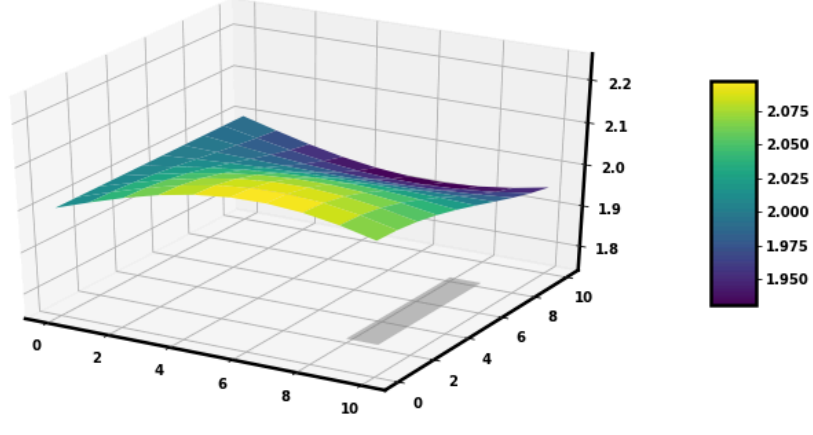
2.2
2.1
2.0
1.9
1.8
0
2
4
6
8
10
2.075
2.050
2.025
2.000
1.975
1.950

The Neural Network is trained ("learns") to mimic the training data

- by solving for the weights $\W_{\llp}$ of each layer $1 \leq\ll\leq L$
- that minimize a loss function

$$\loss = \sum_{i-1}^{m} \loss^{\ip}$$

- where $\loss^{\ip}$ if a function of
    - how much prediction $\hat{\y}^{\ip}$ deviates from true target/label $\y^{\ip}$

The minimization procedure is usually a variant of *Gradient Descent*

The challenge is that the operation of each layer $1 \leq \ll < L$

- is usually not *interpretable*
- we can describe *how* it transforms $\backslash \mathrm{y}_{(\ll -1)}$ to $\backslash \mathrm{y}_{\backslash \mathrm{llp}}$
- but not *why* it is performing the transformation
    - objective (describe) rather than subjection (why)

So the sub-network in the diagram

In [8]:
```
fig_tf_test
```

Out[8]:

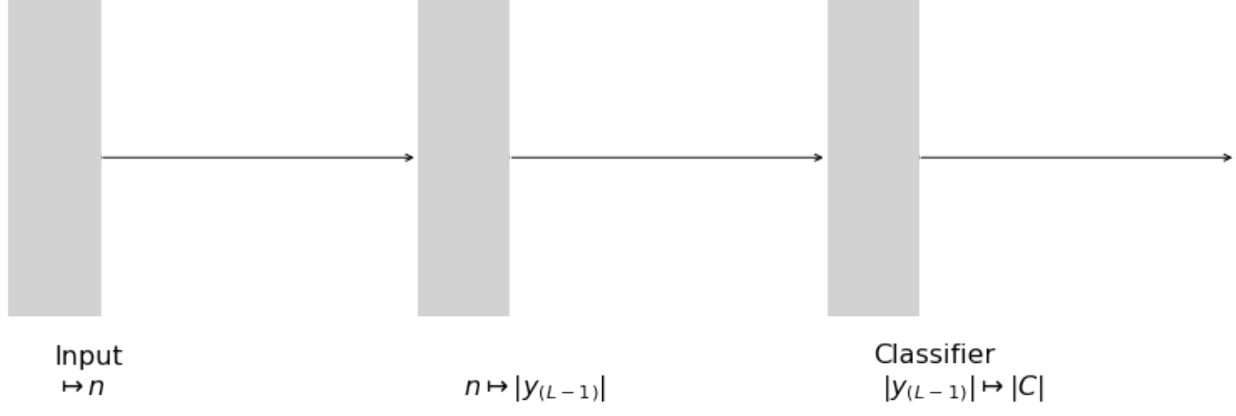

computes an *unknown* function

- from input of length n (the raw input)
- to a vector of length $|\mathrm{y}_{(L-1)}|$

whose purpose is

- to make the final layer L (e.g., the Classifier)
- create predictions (outputs)
- that have a very-low loss loss
    - a good "approximation" of the function described by the training data

In [9]:
```
print("Done")
```

```
Done
```