

Large Margin Classification

So far in the presentation, the difference between the SVC and Logistic Regression classifiers is in the Loss Function.


The SVC is also able to create a "buffer" on either side of the separating boundary.

By making this buffer as wide as possible, an SVC may generalize better.

The buffer is defined by

- Two additional lines
- Parallel to separating boundary
- Same distance (the *margin*) from the separating boundary

```
In [4]: svm_ch = svm_helper.Charts_Helper()  
_ = svm_ch.create_data()  
fig, axs = svm_ch.create_margin(Cs=[10])
```



- The separating boundary is the solid line, whose equation is given in the title
- Each dashed line is
 - Parallel to, and at the same distance from, the separating boundary
 - The distance (measured by length of a line orthogonal to the boundary) from the separating boundary is called the *margin*

The buffer width is twice the margin

In the above plot

- All examples are correctly classified
- There are no examples in the buffer

Requiring these two properties is called *Hard Margin* Classification.



It is somewhat uncommon to be able to achieve the first property (perfect separation of classes).

A more natural Classification task is called *Soft Margin* classification which allows (but penalizes, via the Loss Function) violation of either property.

We re-run the above example with a larger margin

- resulting in the presence of examples in the buffer
 - which were considered as *correctly classified* in the absence of a margin
- which we will consider as *incorrectly classified* in the presence of a non-zero margin
 - and hence will incur a loss

```
In [5]: svm_ch = svm_helper.Charts_Helper()  
_ = svm_ch.create_data()  
fig, axs = svm_ch.create_margin(Cs=[10,.1])
```



We concentrate on Soft Margin Classification going forward.

Achieving a margin

We need to modify the per-example loss to achieve zero loss *only if*

- the example's score is on correct side of the separating boundary
- **and** the example is not in the buffer (i.e., score is exceeds the margin)

This can be achieved by moving the "hinge point" of the Hinge Function

- From 0 to the margin m

This corresponds to a per-example Loss of

$$\mathcal{L}^{(i)} = \max \left(0, m - \dot{\mathbf{y}}^{(i)} * s(\hat{\mathbf{x}}) \right)$$

The above expression achieves zero loss when

$$\hat{s}(\mathbf{x}^{(i)}) \geq m \quad \text{Positive example, } \mathbf{y}^{(i)} = +1$$

$$\hat{s}(\mathbf{x}^{(i)}) \leq -m \quad \text{Negative example, } \mathbf{y}^{(i)} = -1$$

That is:

- an example on the correct side of the separating boundary
- has zero loss
 - only if it is also outside the buffer


How do we choose m ?

As we shall see, a margin $m = 1$ will suffice resulting in

$$\mathcal{L}^{(i)} = \max \left(0, 1 - \mathbf{y}^{(i)} * s(\hat{\mathbf{x}}) \right)$$

Here's the plot

```
In [6]: svmh.plot_hinges(hinge_pt=1)
```



Achieving a large margin

As we observed above, a zero *Classification Loss* occurs when

$$\hat{s}(\mathbf{x}^{(i)}) \geq m \quad \text{Positive example, } \mathbf{y}^{(i)} = +1$$

$$\hat{s}(\mathbf{x}^{(i)}) \leq -m \quad \text{Negative example, } \mathbf{y}^{(i)} = -1$$

The Classification Loss

- penalizes incorrect (and otherwise correct but in the buffer) examples
- but does not force m to be large.

In order to achieve a large margin

- we need to impose a *Margin Penalty* inversely related to the size of m .

We will add this penalty to the Loss Function so that the Loss Function has two terms

- Classification Loss
- Margin Penalty

As previously mentioned:

- there is a simple trick that allows us to consider only $m = 1$

What would happen if we divided both sides of the above inequality (score versus margin) by m ?

- Zero loss occurs when the inequality's right hand side is 1

But dividing both sides by m will affect the parameters Θ used to compute the score

$$\hat{\mathbf{s}} = \Theta_{\text{unscaled}} \cdot \mathbf{x}$$

The unscaled parameters Θ_{unscaled}

- would be rescaled by a factor of $\frac{1}{m}$
- resulting in new parameter values

$$\Theta = \frac{1}{m} \Theta_{\text{unscaled}}$$

Thus

- a large margin
- is associated with *small* Θ
- when dividing both sides of the inequality by m

So we can achieve the *effect* of a large margin

- using constant margin $m = 1$
- by replacing a direct penalty on Margin size
- with a *Regularization Penalty* on parameters size

We enforce the Margin Penalty by the expression

$$\frac{1}{2} \Theta_{-0}^T \cdot \Theta_{-0}$$

as part of the Loss (that is being minimized) in order to force large m

- where Θ_{-0} is a minor variation of Θ as explained below

Notation Our convention is that each example $\mathbf{x}^{(i)}$ has first feature that is the constant 1:

$$\mathbf{x}^{(i)} = [1, \mathbf{x}_1^{(i)}, \dots, \mathbf{x}_n^{(i)}]$$

- Design matrix \mathbf{X} has been augmented with a first column of all 1's
- This allows us to write $\hat{s}(\mathbf{x}^{(i)}) = \Theta^T \cdot \mathbf{x}^{(i)}$
- Θ_0 is the intercept term

Other's (e.g., the Geron book) keep the intercept term *outside* of \mathbf{x}

- Resulting in $\hat{s}(\mathbf{x}^{(i)}) = \Theta^T \cdot \mathbf{x}^{(i)} + \Theta_0$, where \mathbf{x} *does not* have a leading 1
- Geron changes notation from previous chapters (in the "Under the Hood" subsection, page 204)

To avoid confusion, we will write Θ_{-0} to be Θ *excluding* Θ_0

Aside

The mysterious $\frac{1}{2}$ in the Margin Penalty

- Doesn't really affect the overall cost in a significant way
- Will be useful in the mathematical derivations
 - Hint:
 - $\frac{\partial \Theta^2}{\partial \Theta} = 2\Theta$
 - The $\frac{1}{2}$ makes the derivative of the Margin Penalty with respect to Θ exactly Θ
 - The derivative will be used in the optimization of SVM Cost

SVC Loss Function

The final Average Loss Function for the SVC combines

- Classification Loss per-example (penalize incorrect or in-the-buffer predictions)
- Margin Penalty (penalize small margins)

$$\mathcal{L} = \frac{1}{2} \Theta_{-0}^T \cdot \Theta_{-0} + C * \frac{1}{m} \sum_{i=1}^m \max \left(0, 1 - \dot{\mathbf{y}}^{(i)} * s(\hat{\mathbf{x}}^{(i)}) \right)$$

where

$$\hat{s}(\mathbf{x}^{(i)}) = \Theta^T \cdot \mathbf{x}^{(i)}$$

- The first term is the Margin Penalty
- The second term is the average of the per-example losses \mathcal{L}_i
 - weighted by a constant C

What is C ?

- We have two loss terms: Margin Penalty and Average Classification Loss
- C allows us to express a weight for the relative importance of the two loss terms

You should recognize this form of loss function (two loss terms, with relative weight)

- It is like a loss function with a Regularization Penalty

In fact, we will provide a mathematical derivation of the Loss that makes this more apparent.

Let's consider extreme cases of C :

$C = \infty$ No misclassification or buffer violations allowed
forces small margin

$C = 0$ Misclassification and buffer violations unimportant
facilitates larger margin

A high value for C

- May prevent a solution
- Encourage overfitting
 - Less importance on forcing elements of Θ to be zero

A low value for C

- Encourages underfitting
 - More importance on forcing elements of Θ to be zero

In [7]: `print("Done")`

Done

